

Regression_Task_:Saad_Lahlali

November 11, 2021

Contents

1	Settings	2
2	Data Preparation	3
2.1	Reading the data	3
2.2	Standardizing the data	3
2.3	Removing outliers	3
2.4	Exploring the data	4
2.5	Split into training and validation	6
3	Model Selection	7
3.1	Linear Regression	7
3.1.1	Parameters tuning	8
3.1.2	Visualization of LR	10
3.2	Decision Tree Regression	13
3.2.1	Parameters tuning	13
3.3	Gradient-boosted tree regression	15
3.3.1	Parameters tuning	15
3.4	Comparing results	17
3.5	Ensemble Method	18
4	Testing	19
4.1	Reading the data	19
4.2	Making the predictions	19
4.3	Saving the results	19

1 Settings

```
[3]: import pyspark  
sc = pyspark.SparkContext(appName="Regression Task")
```

```
[4]: %matplotlib inline  
  
import matplotlib  
import numpy as np  
import matplotlib.pyplot as plt
```

```
[5]: from pyspark.sql.types import StructType, StructField  
from pyspark.sql.types import DoubleType, IntegerType, StringType  
from pyspark.ml.feature import VectorAssembler, StringIndexer, StandardScaler  
from pyspark.sql import SQLContext
```

2 Data Preparation

2.1 Reading the data

```
[6]: sqlContext = SQLContext(sc)
      #indexer needed so that labels are 0 and 1

      schema = StructType([ StructField("c"+str(i),StringType())]+
                            [ StructField("c"+str(i),DoubleType()) for i in range(1,9)])
      # schema to cast data, can use inferSchema also

      raw_data = sqlContext.read.csv(
          "/content/drive/MyDrive/exo3_train.csv",schema=schema)

      #dropping rows with nulls
      raw_data_nn = raw_data.dropna()

      assembled_data = VectorAssembler(inputCols=["c"+str(i)
          for i in range(1,8)],outputCol="features").
          →transform(raw_data_nn)
```

/usr/local/lib/python3.7/dist-packages/pyspark/sql/context.py:79: FutureWarning:
Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate() instead.
FutureWarning

2.2 Standardizing the data

```
[52]: #standardizing the data
      scale=StandardScaler(inputCol='features',outputCol='standardized_features')
      data_scale=scale.fit(assembled_data)
      data_scale_output=data_scale.transform(assembled_data)
```

2.3 Removing outliers

```
[125]: """
      Applied according to :
      How to remove outliers from multiple columns in pyspark
      using mean and standard deviation
      (from Stackoverflow)
      """

      numeric_cols = ['c1', 'c2', 'c3', 'c4','c5','c6', 'c7']
      mean_std = \
          train \
          .groupBy('c0', 'c8') \
          .agg( \
```

```

    *[f.mean(colName).alias('mean_{}'.format(colName)) for colName in
    ↪numeric_cols],\
    *[f.stddev(colName).alias('stddev_{}'.format(colName)) for colName in
    ↪numeric_cols])

mean_std_min_max = mean_std
for colName in numeric_cols:
    meanCol = 'mean_{}'.format(colName)
    stddevCol = 'stddev_{}'.format(colName)
    minCol = 'min_{}'.format(colName)
    maxCol = 'max_{}'.format(colName)
    mean_std_min_max = mean_std_min_max.withColumn(minCol, f.col(meanCol) - 5 *
    ↪f.col(stddevCol))
    mean_std_min_max = mean_std_min_max.withColumn(maxCol, f.col(meanCol) + 5 *
    ↪f.col(stddevCol))

outliers = train.join(mean_std_min_max, how = 'left', on = ['c0'])
for colName in numeric_cols:
    isOutlierCol = 'is_outlier_{}'.format(colName)
    minCol = 'min_{}'.format(colName)
    maxCol = 'max_{}'.format(colName)
    meanCol = 'mean_{}'.format(colName)
    stddevCol = 'stddev_{}'.format(colName)
    outliers = outliers.withColumn(isOutlierCol,
                                   f.when((f.col(colName) > f.col(maxCol)) |
                                   (f.col(colName) < f.col(minCol)), 1).
                                   otherwise(0))
    outliers = outliers.drop(minCol,maxCol, meanCol, stddevCol)

```

```

[132]: print("With the mean and standard deviation, we have found",
    outliers.filter((f.col("is_outlier_c1")!=f.lit(0)) |
    (f.col("is_outlier_c2")!=f.lit(0)) |
    (f.col("is_outlier_c3")!=f.lit(0)) |
    (f.col("is_outlier_c4")!=f.lit(0)) |
    (f.col("is_outlier_c5")!=f.lit(0)) |
    (f.col("is_outlier_c6")!=f.lit(0)) |
    (f.col("is_outlier_c7")!=f.lit(0))).count(),
    "outliers. (*m)")

```

With the mean and standard deviation, we have found 0 outliers. (*m)

2.4 Exploring the data

```

[8]: import pandas as pd


[9]: sampled_data = data_scale_output.select(["c"+str(i) for i in range(1,9)]).
    sample(False, 0.8).toPandas()
    axs = pd.plotting.scatter_matrix(sampled_data, figsize=(12,12))

```

```

n = len(sampled_data.columns)
for i in range(n):
    v = axs[i, 0]
    v.yaxis.label.set_rotation(0)
    v.yaxis.label.set_ha('right')
    v.set_yticks(())
    h = axs[n-1, i]
    h.xaxis.label.set_rotation(90)
    h.set_xticks(())

```



The most important part of the previous visualization is the last column since it represents the correlation between the target values and the features. Therefore we notice some correlation between some of the variables and the target.

2.5 Split into training and validation

For the model selection, we have to divide our dataset into a training set and a validation set. We set the training ratio to be 80% of the dataset.

```
[10]: train, val = data_scale_output.randomSplit([0.8, 0.2], seed=12345)
```

3 Model Selection

3.1 Linear Regression

```
[11]: from pyspark.ml.regression import LinearRegression
      from pyspark.ml.evaluation import RegressionEvaluator
```

```
[55]: # Building the model
      lr = LinearRegression(featuresCol = 'standardized_features',
                           labelCol='c8',
                           maxIter=10,
                           regParam=0.3,
                           elasticNetParam=0.8)

      lr_model = lr.fit(train)
      print("Coefficients: " + str(lr_model.coefficients))
      print("Intercept: " + str(lr_model.intercept))

      trainingSummary = lr_model.summary
      print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
      print("r2: %f" % trainingSummary.r2)
```

```
Coefficients: [240.9285986458868,292.6629305671314,173.50830089598722,-217.70702
407591259,-114.51404850928375,203.0092587153934,-89.89095612048393]
Intercept: 186.17661860197106
RMSE: 222.477192
r2: 0.719846
```

```
[56]: # Making the predictions
      lr_predictions = lr_model.transform(val)
      lr_predictions.select("prediction","c8","features").show(5)

      lr_evaluator = RegressionEvaluator(predictionCol="prediction", \
                                         labelCol="c8",metricName="r2")

      print("R Squared (R2) on test data = %g" % lr_evaluator.
            ↪evaluate(lr_predictions))
```

```
+-----+-----+-----+
|      prediction|      c8|      features|
+-----+-----+-----+
| 528.4347866392004|1020.0|[446.0,461.0,-27...|
| 361.5438430116119| 354.0|[784.0,-848.0,196...|
| 551.8219666928418| 730.0|[515.0,202.0,-265...|
| 334.6315028322394| 220.0|[408.0,-14.0,-368...|
|471.15826441272804| 736.0|[829.0,-995.0,24...|
+-----+-----+-----+
only showing top 5 rows
```

R Squared (R2) on test data = 0.705601

3.1.1 Parameters tuning

```
[86]: from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit, \
      ↪CrossValidator
      from pyspark.ml.evaluation import RegressionEvaluator

      lr = LinearRegression(featuresCol = 'standardized_features',
                           labelCol='c8',
                           solver="normal")

      # Evaluate model
      lr_evaluator = RegressionEvaluator(predictionCol="prediction",
                                       labelCol="c8",
                                       metricName="rmse")

      # Create ParamGrid for Cross Validation
      lr_paramGrid = (ParamGridBuilder()
                     .addGrid(lr.regParam, [1e-2, 1e-3, 1e-4, 1e-5])
                     .addGrid(lr.elasticNetParam, [0.0, 0.25])
                     .addGrid(lr.maxIter, [1, 2, 5])
                     .build())

      # Create 4-fold CrossValidator
      lr_cv = CrossValidator(estimator = lr,
                           estimatorParamMaps = lr_paramGrid,
                           evaluator = lr_evaluator,
                           numFolds = 4)

      # Run cross validations
      lr_cv_Model = lr_cv.fit(data_scale_output)
      print(lr_cv_Model)

      # Get Model Summary Statistics
      lr_cv_Summary = lr_cv_Model.bestModel.summary
      print("Coefficient Standard Errors: " + str(lr_cv_Summary.
      ↪coefficientStandardErrors))
```

CrossValidatorModel_9a6729a15d9a

Coefficient Standard Errors: [2.6750193208107875, 12.496692273090728,
2.5127277408447344, 1.5465671392467242, 14.840759881676508, 6.373236267906828,
8.546833266216035, 3.8375043042082986]

```
[149]: print("The best model has the following parameters:")
      print("\t"*6, "MaxIter=", lr_cv_Model.bestModel.getMaxIter())
```



```
print("\t"*6,"NetParam=", lr_cv_Model.bestModel.getElasticNetParam())
print("\t"*6,"RegParam=", lr_cv_Model.bestModel.getRegParam())
```

The best model has the following parameters:

```
MaxIter= 1
NetParam= 0.0
RegParam= 0.001
```

```
[61]: # Use validation set here so we can measure the accuracy of our model on new
      ↪data
lr_predictions = lr_cv_Model.transform(val)

lr_score = lr_cv_Summary.rootMeanSquaredError

# cvModel uses the best model found from the Cross Validation
# Evaluate best model
print('RMSE:', lr_score)
```

RMSE: 82.00958198000126

```
[60]: lr_predictions.select("prediction","c8","features").show(10)

lr_evaluator = RegressionEvaluator(predictionCol="prediction", \
                                   labelCol="c8",metricName="r2")

print("R Squared (R2) on test data = %g" % lr_evaluator.
      ↪evaluate(lr_predictions))
```

```
+-----+-----+-----+
|      prediction|      c8|      features|
+-----+-----+-----+
|  927.0563429044515|1020.0|[446.0,461.0,-27...|
|   438.671052424401| 354.0|[784.0,-848.0,196...|
|  811.0902390756099| 730.0|[515.0,202.0,-265...|
|  280.9446575758557| 220.0|[408.0,-14.0,-368...|
|  719.5481187387066| 736.0|[829.0,-995.0,24...|
|   693.980796280758| 631.0|[672.0,-842.0,238...|
| 1037.3125565078024|1179.0|[732.0,222.0,-240...|
|   832.7614185609549| 959.0|[438.0,62.0,-451...|
|  236.28262055046167| 389.0|[322.0,-1484.0,29...|
|-177.21209555850294|-320.0|[200.0,-1283.0,27...|
+-----+-----+-----+

only showing top 10 rows
```

R Squared (R2) on test data = 0.960351

3.1.2 Visualization of LR

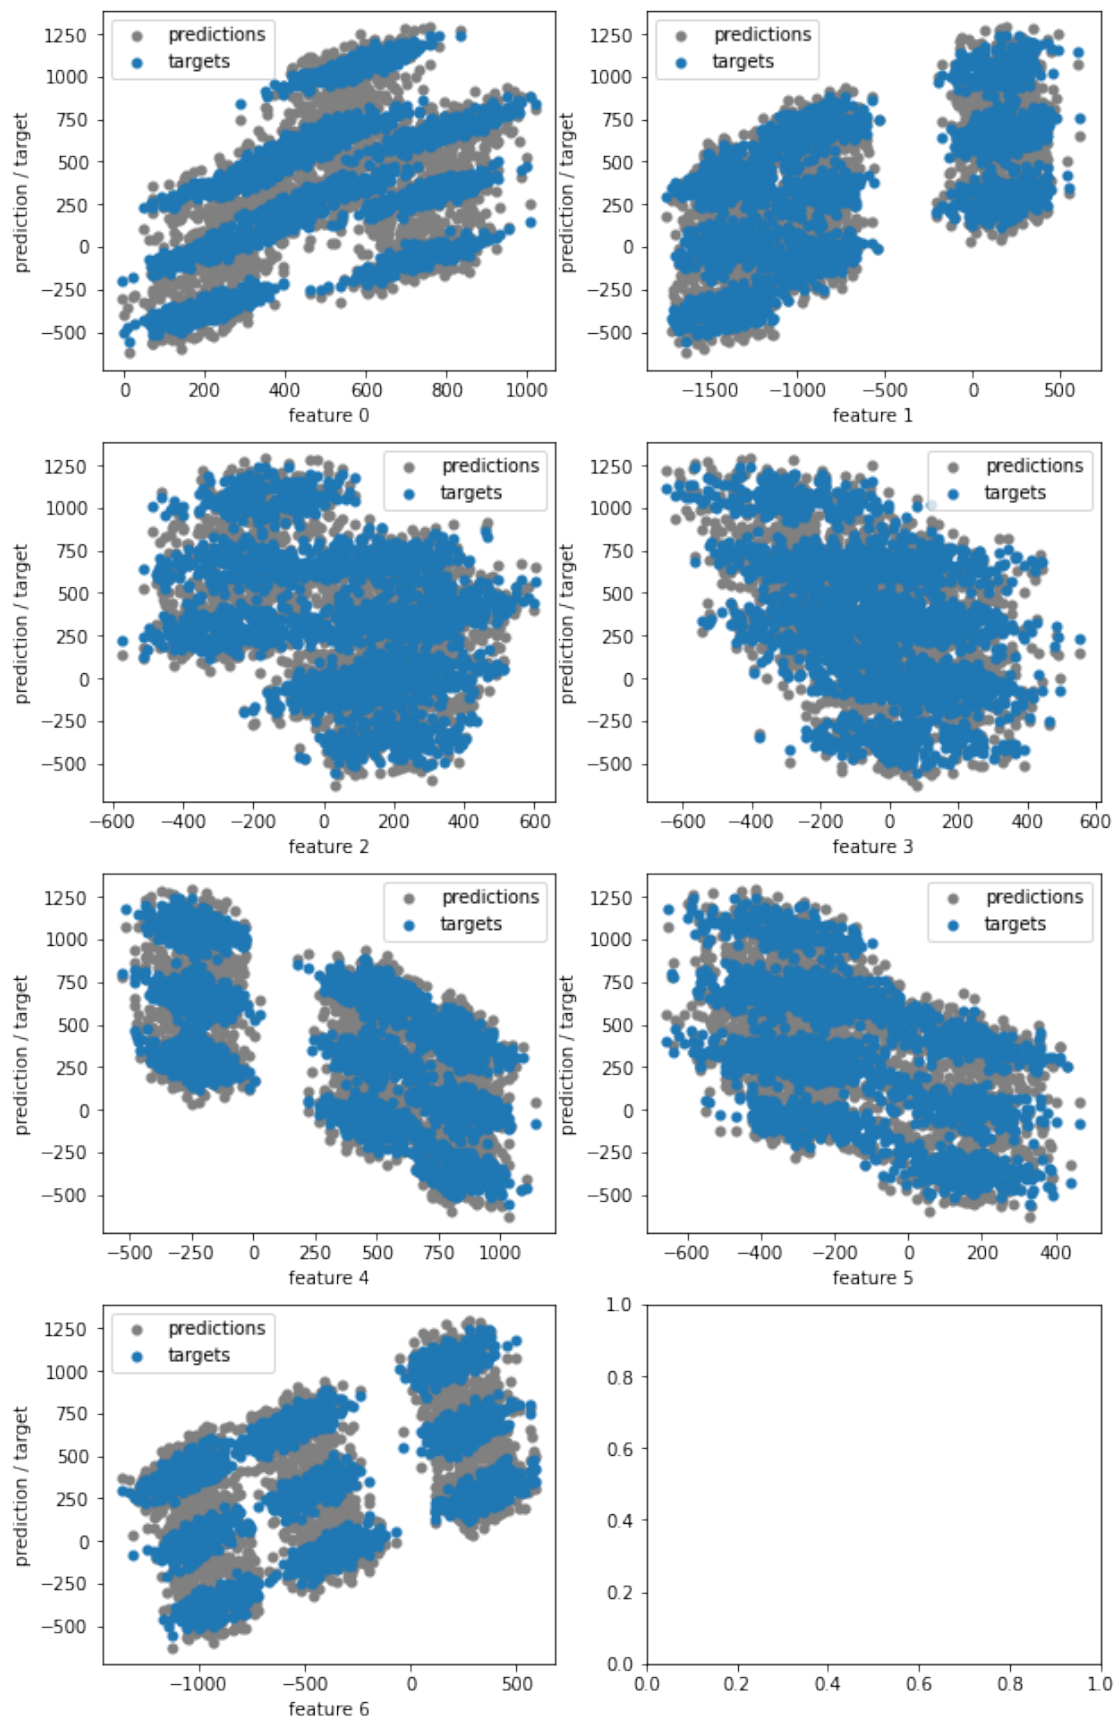
From the following visualization we notice how the target and the predictions superpose in the space which attests of the relative good predictions made by the model.

```
[185]: ft = [w.features for w in lr_predictions.select('features').collect()]
pred = [w.prediction for w in lr_predictions.select('prediction').collect()]
true = [w.c8 for w in lr_predictions.select('c8').collect()]

sort = sorted(range(len(true)), key=lambda k: true[k])
sorted_pred = [pred[i] for i in sort]
sorted_true = [true[i] for i in sort]

fig, ax = plt.subplots(7//2+1, 7%2+1, figsize=(10,17))

for i in range(7):
    c = [f[i] for f in ft]
    sorted_c = [c[i] for i in sort]
    ax[i//2, i%2].scatter(sorted_c, sorted_pred, color='grey', linewidth=0.1,
→label= 'predictions')
    ax[i//2, i%2].scatter(sorted_c, sorted_true, linewidth=0.1, label= 'targets')
    ax[i//2, i%2].set_xlabel('feature ' + str(i))
    ax[i//2, i%2].set_ylabel('prediction / target')
    ax[i//2, i%2].legend()
plt.show()
```



3.2 Decision Tree Regression

```
[ ]: from pyspark.ml.regression import DecisionTreeRegressor
dt = DecisionTreeRegressor(featuresCol='features', labelCol='c8')
dt_model = dt.fit(train)
dt_predictions = dt_model.transform(val)
dt_evaluator = RegressionEvaluator(
    labelCol="c8", predictionCol="prediction", metricName="rmse")
rmse = dt_evaluator.evaluate(dt_predictions)
print("Root Mean Squared Error (RMSE) on validation data = %g" % rmse)
```

Root Mean Squared Error (RMSE) on validation data = 237.499

```
[ ]: dt_model.featureImportances
```

```
[ ]: SparseVector(7, {0: 0.2846, 1: 0.0203, 2: 0.0698, 3: 0.0222, 4: 0.454, 6:
0.1491})
```

3.2.1 Parameters tuning

```
[ ]: # Evaluate model
evaluator = RegressionEvaluator(
    labelCol="c8", predictionCol="prediction", metricName="rmse")

# Create ParamGrid for Cross Validation
paramGrid = ParamGridBuilder().addGrid(dt.maxDepth,
                                         [2, 5, 10, 20, 30]).
                                         addGrid(dt.maxBins,
                                         [10, 40, 100]).
                                         build()

# Create 5-fold CrossValidator
cvs = CrossValidator(estimator=dt,
                    estimatorParamMaps=paramGrid,
                    evaluator=evaluator,
                    # 80% of the data will be used for training, 20% for
→validation.
                    numFolds=5)

# Run cross validations
cvs_model = cvs.fit(train)
print(cvs_Model)
```

CrossValidatorModel_a7f5e34b3bba

```
[ ]: cvs_predictions = cvs_model.transform(val)
evaluator.evaluate(cvs_predictions)
```

```
[ ]: 145.0533679351754
```

```
[ ]: # Use validation set here so we can measure the accuracy of our model on new
    ↪ data
dt_predictions = cvs_model.transform(val)

# cvModel uses the best model found from the Cross Validation
# Evaluate best model
print('RMSE:', evaluator.evaluate(dt_predictions))
```

RMSE: 145.0533679351754

```
[ ]: dt_predictions.select("prediction","c8","features").show(10)

evaluator = RegressionEvaluator(predictionCol="prediction", \
                                labelCol="c8",metricName="r2")

print("R Squared (R2) on test data = %g" % evaluator.evaluate(dt_predictions))
```

```
+-----+-----+-----+
|prediction|    c8|          features|
+-----+-----+-----+
|    1034.0|1020.0|[446.0,461.0,-27...|
|     701.0| 354.0|[784.0,-848.0,196...|
|     741.0| 730.0|[515.0,202.0,-265...|
|     232.0| 220.0|[408.0,-14.0,-368...|
|     739.0| 736.0|[829.0,-995.0,24...|
|     613.0| 631.0|[672.0,-842.0,238...|
|    1168.0|1179.0|[732.0,222.0,-240...|
|     938.0| 959.0|[438.0,62.0,-451...|
|     394.0| 389.0|[322.0,-1484.0,29...|
|    -330.0|-320.0|[200.0,-1283.0,27...|
+-----+-----+-----+
```

only showing top 10 rows

R Squared (R2) on test data = 0.874884

```
[ ]: dt_evaluator = RegressionEvaluator(
    labelCol="c8", predictionCol="prediction", metricName="rmse")
dtr_rmse = dt_evaluator.evaluate(dt_predictions)
print("Root Mean Squared Error (RMSE) on validation data = %g" % rmse)
```

Root Mean Squared Error (RMSE) on validation data = 112.314

3.3 Gradient-boosted tree regression

```
[18]: from pyspark.ml.regression import GBRegressor
      gbt = GBRegressor(featuresCol = 'features', labelCol = 'c8', maxIter=10)
      gbt_model = gbt.fit(train)
      gbt_predictions = gbt_model.transform(val)
      gbt_predictions.select('prediction', 'c8', 'features').show(5)
```

```
+-----+-----+-----+
|      prediction|      c8|      features|
+-----+-----+-----+
| 514.5688085108319|1020.0|[446.0,461.0,-27...|
| 623.8457698169174| 354.0|[784.0,-848.0,196...|
|1010.8209828200846| 730.0|[515.0,202.0,-265...|
| 364.8186941829988| 220.0|[408.0,-14.0,-368...|
| 571.1635744816143| 736.0|[829.0,-995.0,24...|
+-----+-----+-----+
```

only showing top 5 rows

```
[ ]: gbt_evaluator = RegressionEvaluator(
      labelCol="c8", predictionCol="prediction", metricName="rmse")
      rmse = gbt_evaluator.evaluate(gbt_predictions)
      print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

Root Mean Squared Error (RMSE) on test data = 182.743

3.3.1 Parameters tuning

```
[19]: # Evaluate model
      evaluator = RegressionEvaluator(
          labelCol="c8", predictionCol="prediction", metricName="rmse")

      # Create ParamGrid for Cross Validation
      paramGrid = ParamGridBuilder()\
          .addGrid(gbt.maxDepth, [2, 5, 10])\
          .addGrid(gbt.maxIter, [10, 50])\
          .build()

      # Create 5-fold CrossValidator
      gb = CrossValidator(estimator=gbt,
                          estimatorParamMaps=paramGrid,
                          evaluator=evaluator,
                          # 80% of the data will be used for training, 20% for
                          →validation.
                          numFolds=5)
```

```
# Run cross validations  
gb_model = gb.fit(train)  
print(gb_model)
```

CrossValidatorModel_5e8522c33aa8

```
[20]: evaluator = RegressionEvaluator(  
        labelCol="c8", predictionCol="prediction", metricName="rmse")  
gb_predictions = gb_model.bestModel.transform(val)  
gb_score = evaluator.evaluate(gb_predictions)  
print('Rmse obtained is', gb_score)
```

Rmse obtained is 127.95867390709157

3.4 Comparing results

```
[ ]: I = ['Linear Regression', 'Decision Tree', 'Gradient Boost']
L = [lr_score, dtr_rmse, gb_score]

labels = L
x = np.arange(len(labels))
width = 0.35

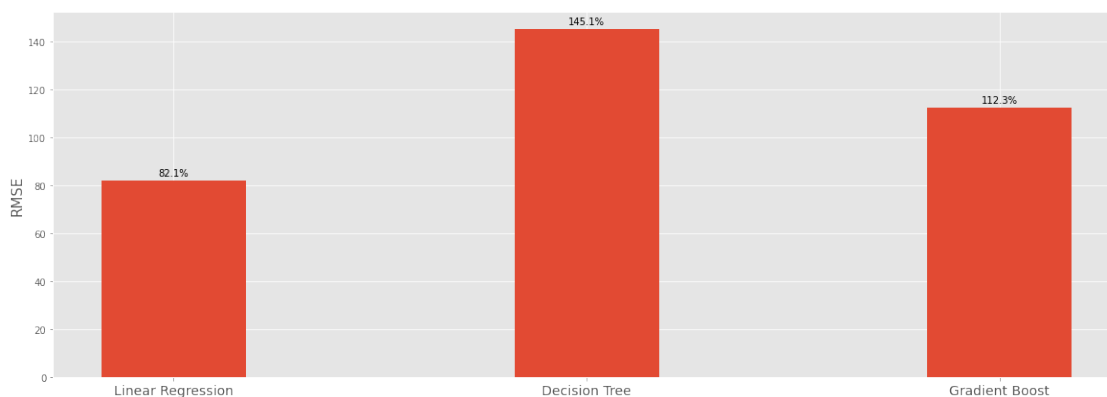
fig, ax = plt.subplots(figsize=(20,7))
rects1 = ax.bar(x, L, width)

def autolabel(rects, labels):
    i=0
    for rect in rects:
        height = rect.get_height()
        ax.annotate(s = str(labels[i])+'%',
                    xy=(rect.get_x() + rect.get_width() / 2, height),
                    xytext=(0, 3), # 3 points vertical offset
                    textcoords="offset points",
                    ha='center', va='bottom')
        i+=1

labels1 = [round(A_SVM_sorter, 1) for A_SVM_sorter in L]
autolabel(rects1, labels1)

ax.set_ylabel('RMSE', size = 15)
ax.set_xticks(x)
ax.set_xticklabels(I, size = 14)

plt.show()
```



We get much better results on the Linear Regression than in the two other models. Therefore we will choose this model for the rest of the project.

3.5 Ensemble Method

```
[36]: lr_df = lr_predictions.select(lr_predictions["c0"].alias("c0_lr"),  
    ↳lr_predictions["prediction"].alias("lr_pred"), lr_predictions["c8"])  
gb_df = gb_predictions.select(gb_predictions["c0"].alias("c0_gb"),  
    ↳gb_predictions["prediction"].alias("gb_pred"))  
  
[142]: from pyspark.sql.functions import col  
    # finding good compromise between weight of prediction to Gradient Boost and  
    ↳Linear Regression  
    for lr_a in range(0,15):  
  
        marksColumns = [col('lr_pred')] + [col('gb_pred')]*lr_a  
  
        averageFunc = sum(x for x in marksColumns)/len(marksColumns)  
  
        df = gb_df.join(lr_df, lr_df["c0_lr"]==gb_df["c0_gb"], "inner")  
  
        ens = df.withColumn('ens_pred', averageFunc).select(col("c0_gb").  
    ↳alias("index"), col("ens_pred").alias("prediction"), col("c8").  
    ↳alias("label"))  
  
        evaluator = RegressionEvaluator(  
            labelCol="label", predictionCol="prediction", metricName="rmse")  
        score = evaluator.evaluate(ens)  
        print(lr_a, score)
```

```
0 222.50513985032347  
1 148.66018770825247  
2 133.34879846641687  
3 128.55487052280907  
4 126.73643162225233  
5 125.98846107999964  
6 125.68639417619252  
7 125.5878893870486  
8 125.58736711333606  
9 125.63489774271443  
10 125.70545467580429  
11 125.78598818210368  
12 125.86953638095635  
13 125.9523657479568  
14 126.03250974003987
```

The ensemble method path was not conclusive, the results were worse than for a simple Linear Regression

4 Testing

4.1 Reading the data

```
[ ]: schema = StructType([ StructField("c"+str(i),StringType())]+[  
    ↳StructField("c"+str(i),DoubleType()) for i in range(1,9)])  
# schema to cast data, can use inferSchema also  
  
raw_data = sqlContext.read.csv("/content/drive/MyDrive/exo3_predict.  
    ↳csv",schema=schema)  
  
assembled_data = VectorAssembler(inputCols=["c"+str(i) for i in_  
    ↳range(1,8)],outputCol="features").transform(raw_data)  
  
#standardizing the data  
scale=StandardScaler(inputCol='features',outputCol='standardized_features')  
data_scale=scale.fit(assembled_data)  
data_scale_output=data_scale.transform(assembled_data)
```

4.2 Making the predictions

```
[ ]: test_lr_predictions = lr_cv_Model.transform(data_scale_output)
```

4.3 Saving the results

```
[ ]: pred = [w.prediction for w in test_lr_predictions.select('prediction').  
    ↳collect()]  
ind = [w.c0 for w in test_lr_predictions.select('c0').collect()]  
  
with open('exo3.csv','wb') as file:  
    for p, i in zip(pred, ind):  
        file.write((str(i)+','+str(p)+'\n').encode())
```